

**Lecture 36 [15.05.2020]**

## **NoC Router Microarchitecture**




**John Jose**

**Assistant Professor**

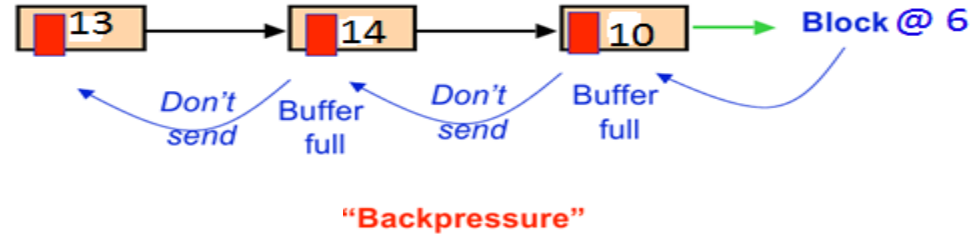
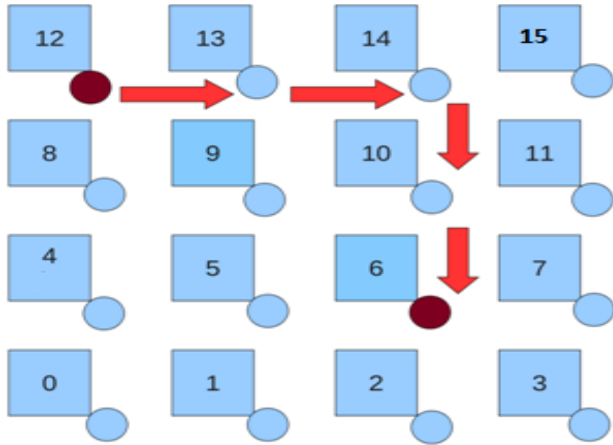
**Department of Computer Science & Engineering  
Indian Institute of Technology Guwahati, Assam.**

# Building Blocks of NoC

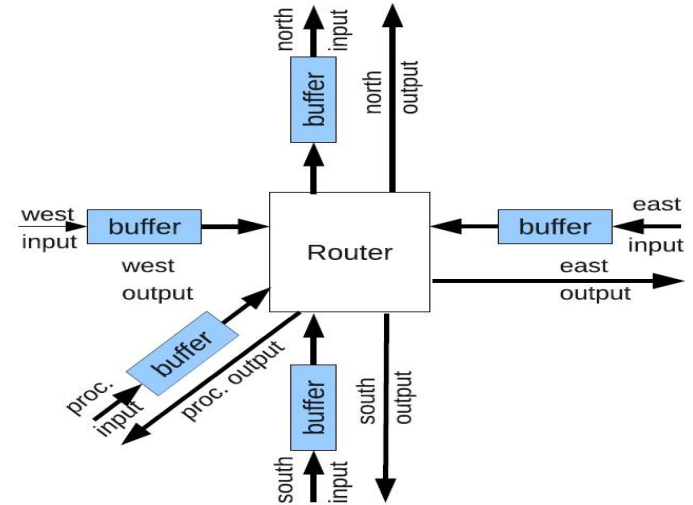


- ❖ **Topology**
- ❖ **Routing**
- ❖ **Flow control**
- ❖ **Router micro-architecture**

# Flow Control




- ❖ Upstream router should know the buffer availability of downstream router.
- ❖ Credit should be exchanged between routers by handshake signals




# Store and Forward Flow Control

- ❖ Store and Forward Packet based flow control
  - ❖ Packet copied entirely into network router before moving to the next node
  - ❖ Flow control unit is the entire packet
- ❖ Leads to high per-packet latency
- ❖ Requires buffering for entire packet in each node




# Wormhole Flow Control



- ❖ Packets broken into smaller flits
- ❖ Flits are sent across the fabric in a *wormhole fashion*
  - ❖ Body follows head, tail follows body
  - ❖ Pipelined
  - ❖ If head blocked, rest of packet stops
  - ❖ Routing (src/dest) information only in head
- ❖ Lower latency, efficient buffer utilization
- ❖ Occupies resources across multiple routers

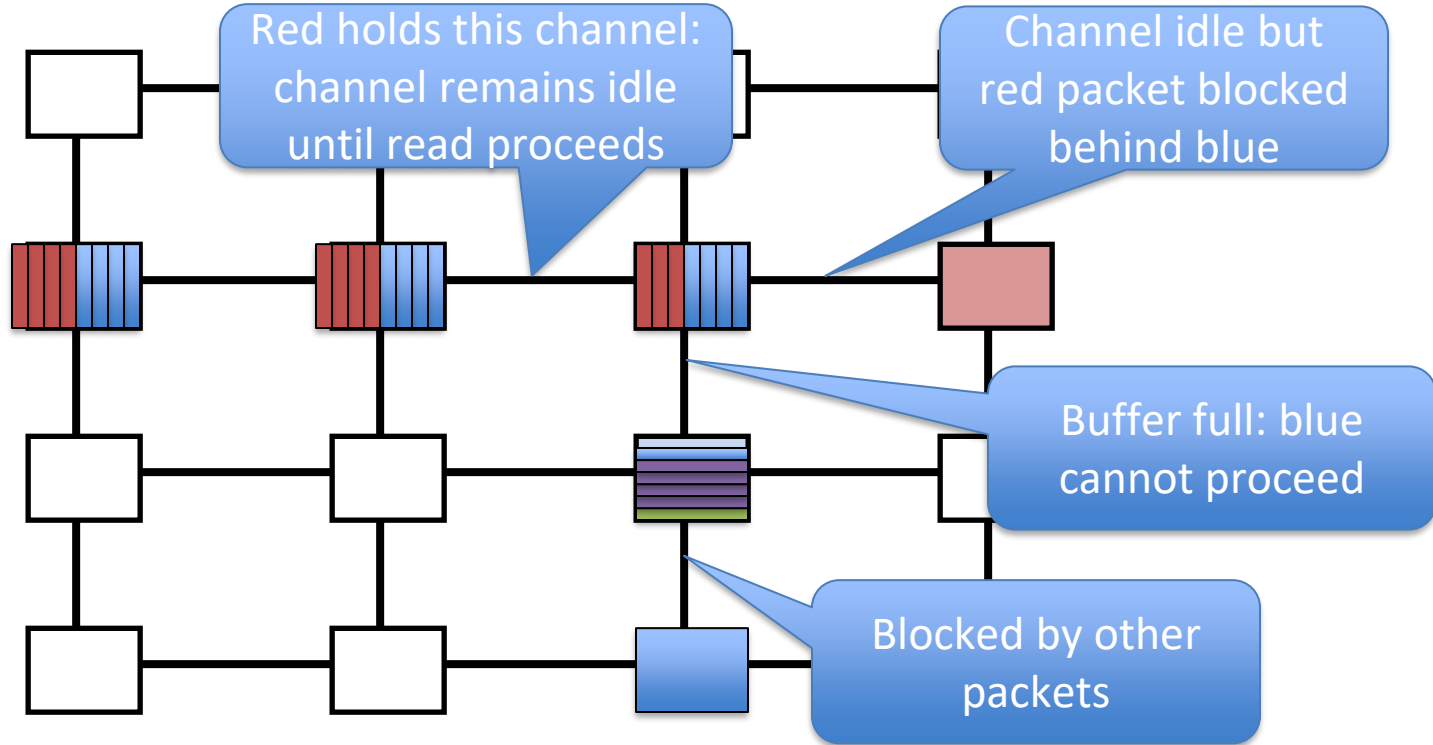
# Head of Line Blocking

- ❖ Suffers from **head of line blocking**
  - ❖ If head flit cannot move due to contention, another worm cannot proceed even though links may be idle




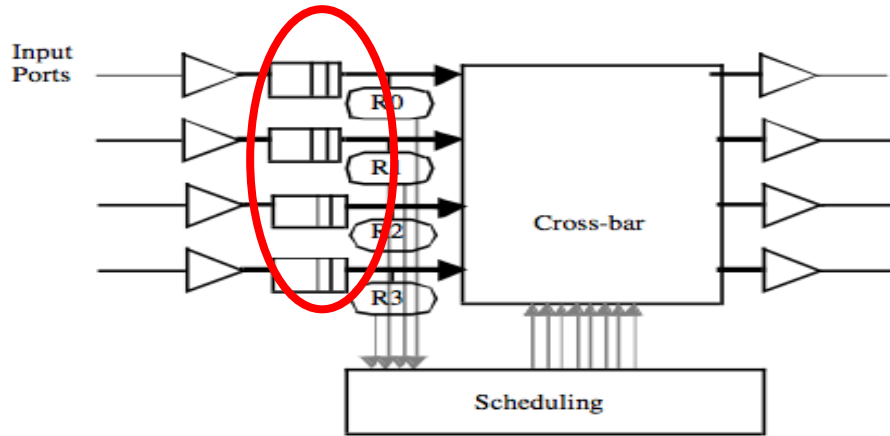
**HOL Blocking**

# Head of Line Blocking




# Virtual Channel Flow Control

- ❖ Multiplex multiple channels over one physical channel
- ❖ FIFO buffers replaced with multilane buffers
- ❖ Divide up the input buffer into multiple buffers sharing a single physical channel



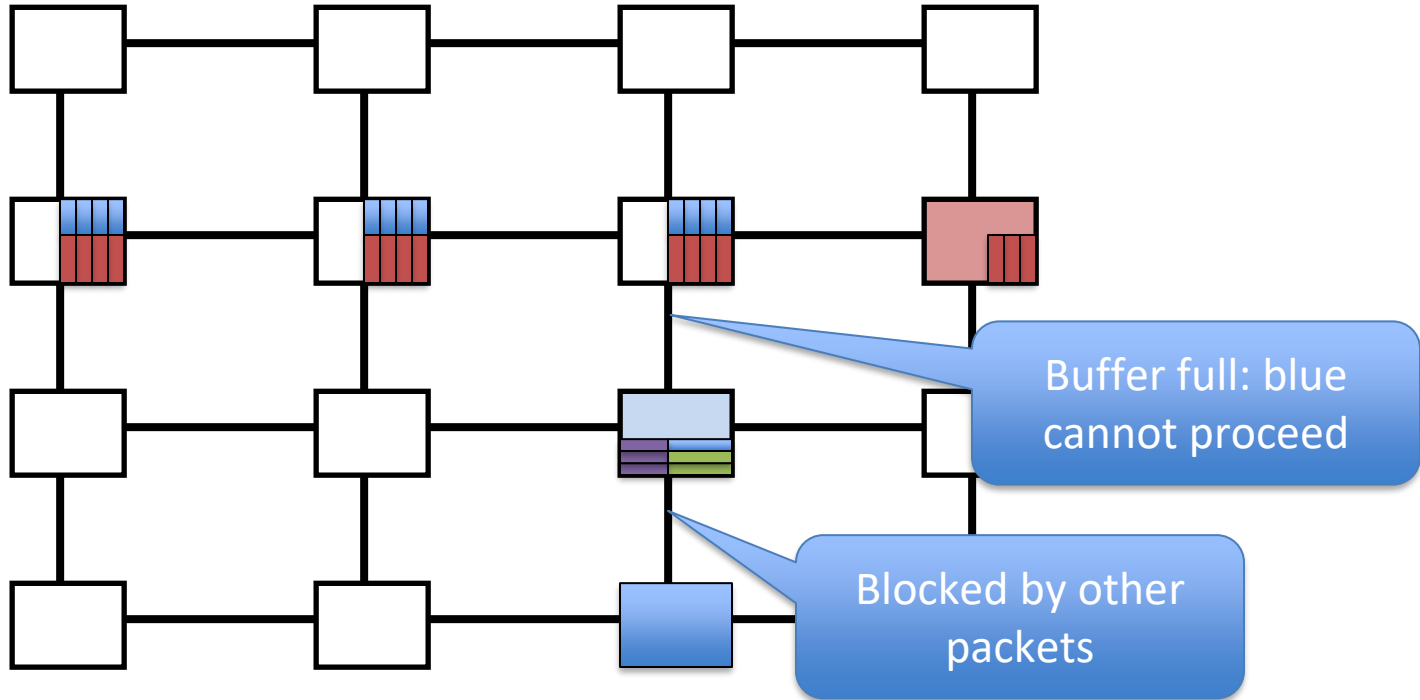


# Virtual Channel Flow Control

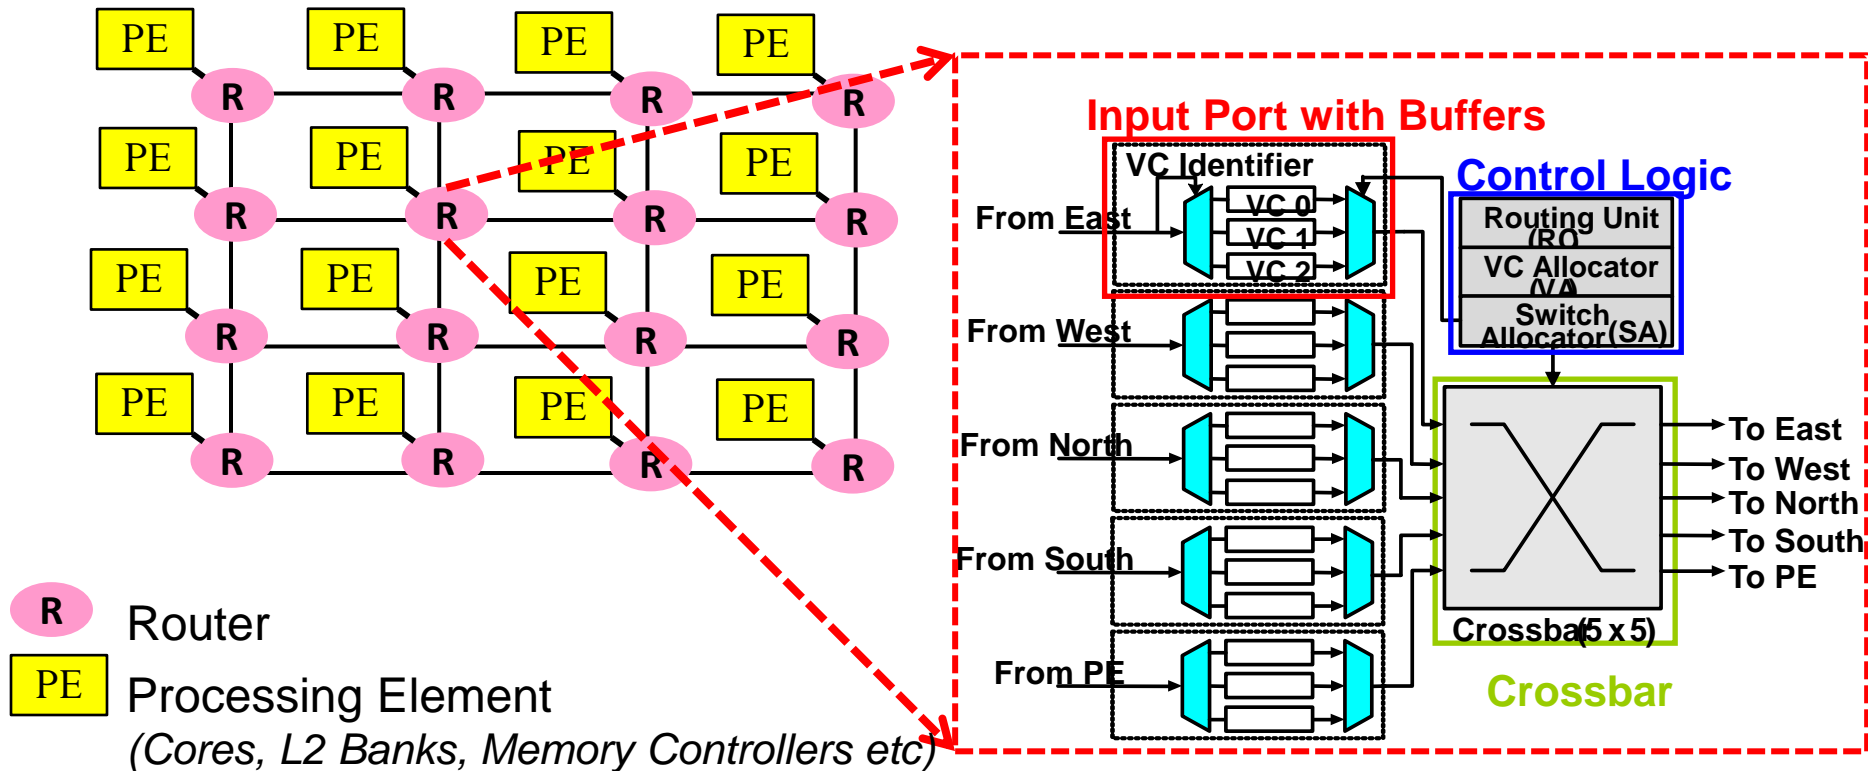


- ❖ VCs are allocated once at each router to the head flit and remaining flits of the packet inherit the same VC
- ❖ Flits of different packets can be interleaved on the same physical channel
- ❖ VCs avoid deadlocks

# Virtual Channel Flow Control

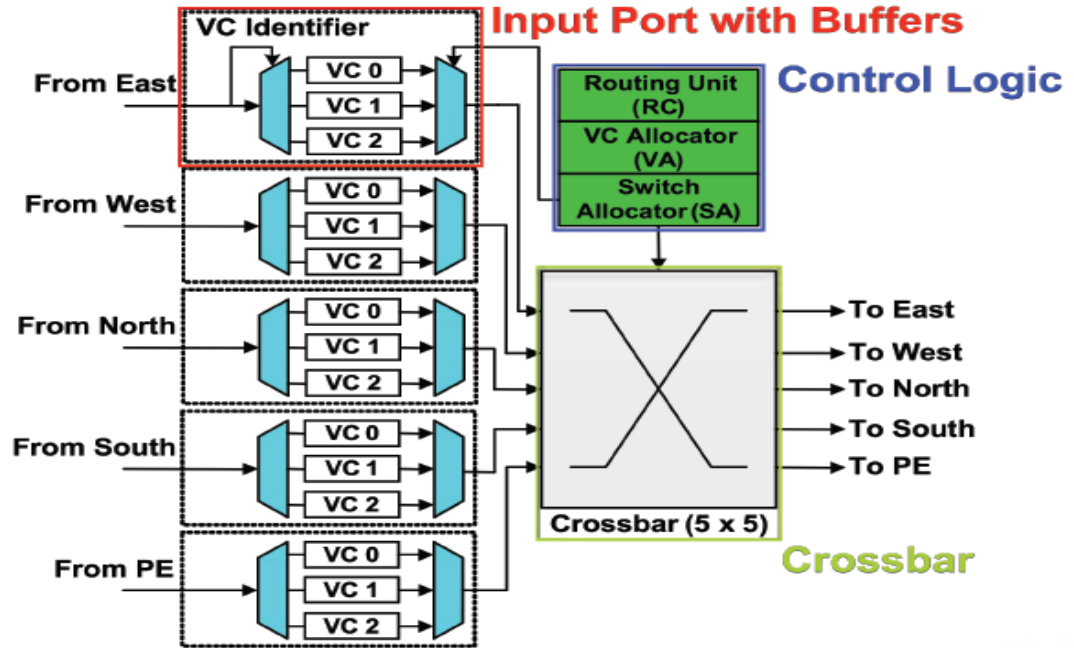


# Input Buffered NoC Router




# Functions of a Router

- ❖ Buffering of flits
- ❖ Route computation
- ❖ VC allocation
- ❖ Switch Allocation
- ❖ Switch Traversal
- ❖ Link Traversal



# Router Pipeline



❖ Five logical stages

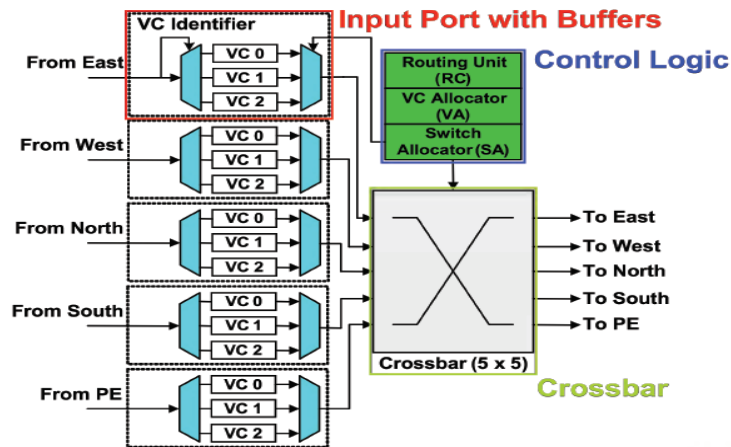
BW: Buffer Write

❖ RC: Route computation


❖ VA: Virtual Channel Allocation

❖ SA: Switch Allocation

❖ ST: Switch Traversal




# Wormhole Router Timeline



- ❖ Route computation performed once per packet
- ❖ Virtual channel allocated once per packet
- ❖ Body and tail flits inherit this information from head flit

# Dependencies in a Router



- ❖ Dependence between output of one module and input of another
  - ❖ Determine critical path through router
  - ❖ Cannot bid for switch port until routing performed

# Lookahead Routing

- ❖ At current router perform routing computation for next router

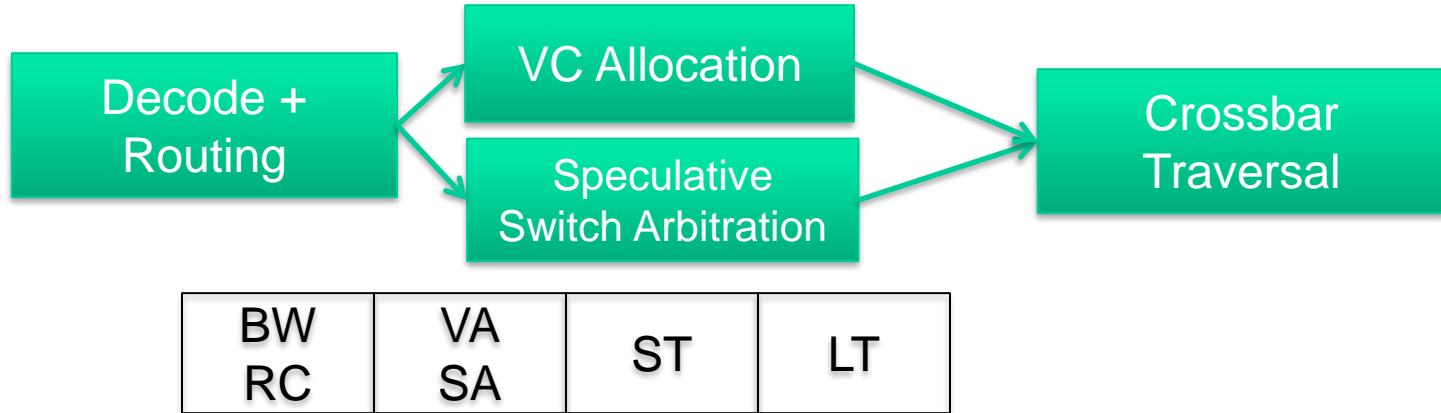
- ❖ Overlap with BW

BW RC	VA	SA	ST	LT
----------	----	----	----	----

- ❖ Pre-computing route allows flits to compete for VCs immediately after BW
- ❖ RC decodes route header
- ❖ Routing computation needed at next hop
  - ❖ Can be computed in parallel with VA





# Speculative Routing




- ❖ Assume that Virtual Channel Allocation stage will be successful
  - ❖ Valid under low to moderate loads
- ❖ Entire VA and SA in parallel
- ❖ If VA unsuccessful (no virtual channel returned)
  - ❖ Must repeat VA/SA in next cycle

# Selection Strategy


When there are multiple possible paths for a packet at a router, which one to choose ?




# Input / Output Channel Selection



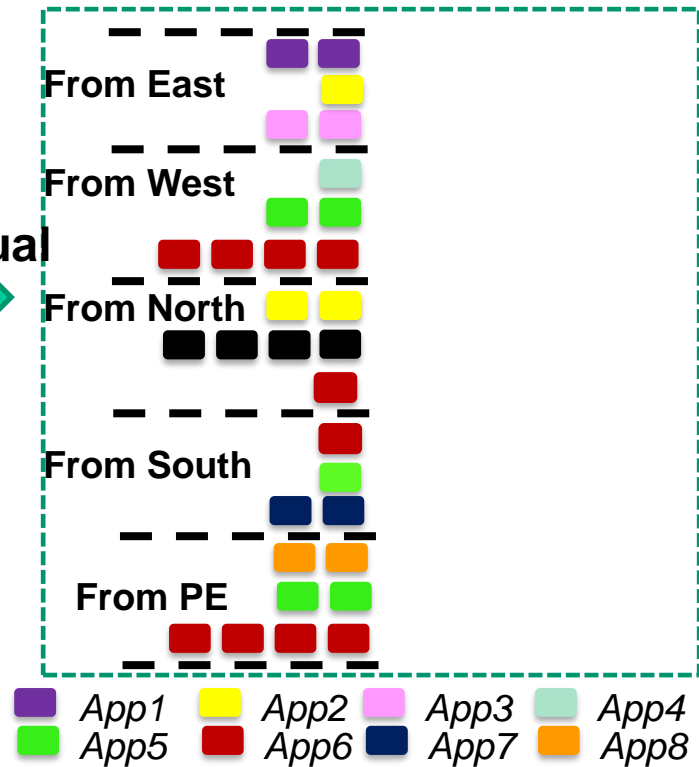
## Output Channel Selection Input Channel Selection



# Switch Level Packet Scheduling in NoC



Conceptual  
View




# Application-to-Core Mapping Policies

Applications



Cores



How to map applications to cores?

# Application-to-Core Mapping Policies

## ❖ Application To Core Mapping

❖ Clustering

❖ Balancing

❖ Isolation

❖ Radial mapping




# Task Scheduling

- ❖ Traditional
  - ❖ When to schedule a task? – Temporal
- ❖ Many-Core
  - ❖ When to schedule a task? – Temporal
  - ❖ **Where to schedule a task? – Spatial**
- ❖ Spatial scheduling impacts performance of memory hierarchy
- ❖ Latency is impacted by interference in NoC, memory, and caches

# On-Chip Communication

Application




**Memory  
Controller**



**Shared  
Cache Bank**




# Clustering



▲ Memory  
Controller

**Inefficient data mapping to memory and caches**

# Clustering




**Improved Locality**

**Reduced Interference**

# Clustering

- ❖ Locality aware page replacement policy
- ❖ When allocating free page, give preference to pages belonging to the cluster's memory controllers (MCs)



# Balancing

## Applications




Heavy



Light

## Cores



**Too much load in clusters with heavy applications**

# Balancing

# Applications




# Heavy



## Light

## Cores



# Isolation

## Applications



Sensitive



Light




Medium



Heavy

## Cores



Isolate **sensitive** applications to a cluster

Balance load for remaining applications across clusters

# Isolation

- ❖ How to estimate sensitivity?
  - ❖ High Miss— high misses per kilo instruction (MPKI)
  - ❖ Low MLP— high relative stall cycles per miss (STPM)
  - ❖ Sensitive if  $MPKI > \text{Threshold}$  & relative STPM is high
- ❖ Whether to or not to allocate cluster to sensitive applications?

# Radial Mapping

Applications



Sensitive



Light




Medium



Heavy

Cores



Map applications **that benefit most from** being close to **memory controllers** close to these resources





**johnjose@iitg.ac.in**  
**<http://www.iitg.ac.in/johnjose/>**