

Database Management Systems

Vijaya Saradhi

IIT Guwahati

Mon, 17th Feb 2020

DDL

Overview

- Used in **creating** tables that is entities, relations
- Defining domains for each column that is attribute
- Express constraints on tables
- Modify tables
- Modify constraints
- Delete tables, columns within tables and constraints
- User need to have privileges for performing these operations
- Typically database administrator perform these operations
- Database users perform DML

DDL - Create

Permanent table

```
CREATE TABLE student(  
    roll_number CHAR(20),  
    name CHAR(30),  
    login CHAR(20),  
    age INT);
```

DDL - Create temporary table

Temporary table

```
CREATE TEMPORARY TABLE student(  
    roll_number CHAR(20),  
    name CHAR(30),  
    login CHAR(20),  
    age INT);
```

DDL - Expressing - Keys

Expressing Keys

```
CREATE TABLE student(  
    roll_number CHAR(20),  
    name CHAR(30),  
    login CHAR(20),  
    age INT,  
    UNIQUE(login),  
    PRIMARY KEY(roll_number));
```

Referential Integrity Constraint

Deletion

- Deletion of student with id 53666 not only affects student table but also grades table
- Possible scenarios of deletion
 - Delete all from grade which references 53666
 - Disallow the deletion of the student row 53666
 - Set the `student_id` to some default value
 - Set the `student_id` to `null` value.
 - However primary keys cannot assume null values and hence this is not a feasible option

DDL - Expressing Foreign Keys - 01

```
CREATE TABLE student(  
    roll_number CHAR(20),  
    name CHAR(30) NOT NULL,  
    login CHAR(20),  
    age INT NOT NULL,  
    UNIQUE(login),  
    PRIMARY KEY(roll_number));
```

DDL - Expressing Foreign Keys - 02

```
CREATE TABLE course(  
  cid CHAR(6),  
  title CHAR(20) NOT NULL,  
  credits INT NOT NULL,  
  PRIMARY KEY(cid));
```


DDL - Expressing Foreign Keys - 03

```
CREATE TABLE registers(  
    rn CHAR(20),  
    course_id CHAR(6),  
    PRIMARY KEY(rn, course_id),  
    FOREIGN KEY(rn) REFERENCES student(roll_number),  
    FOREIGN KEY(course_id) REFERENCES course(cid)  
);
```

DDL - Expressing Foreign Keys - 04

Specify actions

- What happens when a student row gets deleted?
- What happens when a student row gets updated?
- What happens when a course row gets deleted?
- What happens when a course row gets updated?

```
CREATE TABLE registers(  
    rn CHAR(20),  
    course_id CHAR(6),  
    PRIMARY KEY(rn, course_id),  
    FOREIGN KEY(rn) REFERENCES student(roll_number)  
    ON DELETE CASCADE ON UPDATE NO ACTION,  
    FOREIGN KEY(course_id) REFERENCES course(cid)  
    ON DELETE CASCADE ON UPDATE CASCADE  
);
```

DDL - Expressing Foreign Keys - 05

```
CREATE TABLE manager(  
    supervisor_id CHAR(20),  
    supervisee_id CHAR(20),  
    PRIMARY KEY(supervisor_id, supervisee_id),  
    FOREIGN KEY(supervisor_id) REFERENCES employee(eid)  
    FOREIGN KEY(supervisee_id) REFERENCES employee(eid)  
);
```

DDL - Foreign Keys - Scenario - 1

Delete all from grade which registers 53666

```
CREATE TABLE registers(  
    rn CHAR(20),  
    course_id CHAR(6),  
    PRIMARY KEY(rn, course_id),  
    FOREIGN KEY(rn) REFERENCES student(roll_number)  
    ON DELETE CASCADE  
);
```

DDL - Foreign Keys - Scenario - 2

Disallow the deletion of the student row

```
CREATE TABLE registers(  
    rn CHAR(20),  
    course_id CHAR(6),  
    PRIMARY KEY(rn, course_id),  
    FOREIGN KEY(rn) REFERENCES student(roll_number)  
    ON DELETE RESTRICT  
);
```

DDL - Foreign Keys - Scenario - 3

Set the student_id to some default value

```
CREATE TABLE registers(  
    rn CHAR(20),  
    course_id CHAR(6),  
    PRIMARY KEY(rn, course_id),  
    FOREIGN KEY(rn) REFERENCES student(roll_number)  
    ON DELETE SET DEFAULT rn='1234'  
);
```

DDL - Foreign Keys - Scenario - 4

Set the student_id to NULL value

```
CREATE TABLE registers(  
    rn CHAR(20),  
    course_id CHAR(6),  
    PRIMARY KEY(rn, course_id),  
    FOREIGN KEY(rn) REFERENCES student(roll_number)  
    ON DELETE SET DEFAULT rn=NULL  
);
```

DDL - Expressing Foreign Keys - 06

- A constraint is checked at the **end** of every SQL statement
- Checks for constraint violations
- SQL statements gets rejected in the case of constraint violations
- Some times this causes inflexibility

DDL - Expressing Foreign Keys - 06a

Table 1

```
CREATE TABLE student(  
    roll_number CHAR(20),  
    name CHAR(30),  
    login CHAR(20),  
    age INT,  
    honors CHAR(10) NOT NULL,  
    UNIQUE(login),  
    PRIMARY KEY(roll_number),  
    FOREIGN KEY (honors) REFERENCES courses(cid)  
);
```

DDL - Expressing Foreign Keys - 06b

Table 2

```
CREATE TABLE course(  
  cid CHAR(6),  
  title CHAR(20) NOT NULL,  
  credits INT NOT NULL,  
  grader CHAR(20) NOT NULL,  
  PRIMARY KEY(cid),  
  FOREIGN KEY(grader) REFERENCES student(roll_number)  
);
```

DDL - Expressing Foreign Keys - 06c

Deffer constraint

Disable foreign key checks

```
SET foreign_key_checks = 0;
```

DDL - Expressing Foreign Keys - 06d

Enable constraint check

Enable foreign key checks

```
SET foreign_key_checks = 1;
```

DDL - Default Constraint

Setting Default Values

```
CREATE TABLE student(  
    sid char(9) PRIMARY KEY,  
    name varchar(30),  
    phone char(10) DEFAULT '1234567890'  
);
```

DDL - Create Domain

Domain Constraints

- The DOMAIN is a new schema element
- You can think of this as providing alias to the an SQL data type statement
- Allows to declare in **in-line macro**
- Syntax: CREATE DOMAIN <domain name> AS < data type >

```
1 CREATE DOMAIN CPI_DATA AS REAL CHECK (value >= 0 AND
   value <= 10);
2
3 CREATE TABLE student(
4     sid char(9) PRIMARY KEY,
5     name varchar(30),
6     cpi CPI_DATA
7 );
8
```

DDL - Naming Constraint

Naming constraints

- Every constraint can be given a name
- Names are useful in creating, modifying and deleting constraints on tables
- Every constraint is prefixed with syntax **CONSTRAINT [symbol]** followed by the actual constraint
- In the following example, c1, c2 and c3 are the names given to each of the constraint

```
CONSTRAINT c1 UNIQUE(login)
CONSTRAINT c2 PRIMARY KEY(roll_number)
CONSTRAINT c3 FOREIGN KEY (honors) REFERENCES courses(cid)
```

DDL - Creating Indexes

Indexes

- Provide handle on adding indexes to existing tables
- Following example creates an index using the first 10 characters of the name column

```
CREATE INDEX part_of_student_name ON student(name(10));
```


DDL - Adding a column

Altering Table

R			
c1	c2	c4	c5

- Adding a column between c2 and c4

```
1 ALTER TABLE R ADD COLUMN c3 INT AFTER c2 ;  
2
```

R				
c1	c2	c3	c4	c5

DDL - Adding a column at the beginning

Altering Table

R			
c2	c3	c4	c5

- Adding a column c1 at the beginning

```
1 ALTER TABLE R ADD COLUMN c1 INT FIRST ;  
2
```

R				
c1	c2	c3	c4	c5

DDL - Adding a column at the end

Altering Table

R			
c1	c2	c3	c4

- Adding a column c1 at the end

```
1 ALTER TABLE R ADD COLUMN c5 INT;  
2
```

R				
c1	c2	c3	c4	c5

DDL - Dropping a column

Altering Table

R				
c1	c2	c3	c4	c5

- Dropping the column c1

```
1 ALTER TABLE R DROP COLUMN c1 ;  
2
```

R			
c2	c3	c4	c5

DDL - Adding Constraints

Primary Key

```
CREATE TABLE R(c1 INT, c2 INT, c3 INT, c4 INT);
```

R			
c1	c2	c3	c4

- Adding a primary key c1

```
1 ALTER TABLE R ADD CONSTRAINT my_c1 PRIMARY KEY(c1);
```

```
2
```

DDL - Adding Constraints

Foreign Key

```
CREATE TABLE R(c1 INT, c2 INT, c3 INT, c4 INT, PRIMARY KEY(c1
));
CREATE TABLE S(s1 INT, s2 INT, PRIMARY KEY(s1);
```

- Adding a primary key c2 to R

```
1 ALTER TABLE R ADD CONSTRAINT my_c2_fkey FOREIGN KEY(c2)
  REFERENCES S(s1);
```

2

DDL - Dropping Constraints

Primary Key

```
ALTER TABLE R DROP CONSTRAINT my_c1;
```

Foreign Key

```
ALTER TABLE R DROP CONSTRAINT my_c2_fkey;
```

DDL - Changing Domains

Altering Attribute Domains

```
ALTER TABLE R CHANGE c3 c3 CHAR(20);
```

```
ALTER TABLE R CHANGE c3 new_c3 CHAR(20);
```

One has to be careful while changing the domains when with columns are either primary key or foreign key constraints.

DDL - Default Constraint

Expressing Default Constraint

```
CREATE TABLE R(c1 INT, c2 INT DEFAULT 441, PRIMARY KEY(c1))
```

Primary key vs temporal key

Example Schema

- eid and pcn stand for primary key
- Only in the absence of timed attributes
- start_date and end_date are included in the relation
- No employee can have a particular position twice at the same time.
- eid, pcn, start_date, end_date not a primary key

eid	pcn	start_date	end_date
123	900225	01-Jan-1996	01-June-1996
123	900225	01-Apr-1996	01-Oct-1996

Primary key vs temporal key

```
CREATE TABLE Incumbents( eid INT, pcn INT, start_date date,
end_date date,
CHECK(
    NOT EXISTS (
        SELECT *
        FROM Incumbents as l1
        WHERE 1 <
            (SELECT COUNT(eid)
             FROM Incumbents as l2
             WHERE l1.eid = l2.eid
             AND l1.pcn = l2.pcn
             AND l1.start_date < l2.end_date
             AND l2.start_date < l1.end_date)
        )
    AND NOT EXISTS (
        SELECT *
        FROM Incumbents AS l1
        WHERE l1.eid is null OR l1.pcn is null
    )
)
```

SELECT

Overview

- Consists of **SIX** clauses
- Combines **selection** and **projection** operators
- Optionally the following are specified
 - Extended operations
 - Groupy
 - sort (order by)

SELECT list of attributes

FROM list of tables

WHERE Condition

GROUP BY list of attributes

HAVING CONDITION

ORDER BY list of attributes

Algebraic Operators and SQL

Overview

σ, π SELECT, FROM, WHERE

× comma separated table list after FROM clause

× table_1 CROSS JOIN table_2

⋈ table_1 JOIN table_2

Theta Join table_1 JOIN table_2 ON Condition

Re-naming AS: SELECT bname AS boat_name FROM Boats

Algebraic Operators and SQL

Operators

- ∪ UNION
- ∩ INTERSECTION (not available in all DBs)
- EXCEPT (not available in all DBs)

Selection

 $\sigma_{attr3 \geq 6}(table1)$

table1		
attr1	attr2	attr3
1	2	5
3	4	6
1	2	7
1	2	7

```
SELECT attr1 , attr2 , attr3
FROM table1
WHERE attr3 >= 6;
```

Selection

 $\sigma_{attr3 \geq 6}(table1)$

table1		
attr1	attr2	attr3
1	2	5
3	4	6
1	2	7
1	2	7

```
SELECT *  
FROM   table1  
WHERE  attr3 >= 6;
```


Projection

 $\pi_{attr1, attr2}(table1)$

table1		
attr1	attr2	attr3
1	2	5
3	4	6
1	2	7
1	2	7

```
SELECT attr1 , attr2
FROM table1;
```

Projection

 $\pi_{attr3}(table1)$

table1		
attr1	attr2	attr3
1	2	5
3	4	6
1	2	7
1	2	7

```
SELECT attr3  
FROM table1;
```

Selection AND Projection

$$\pi_{attr2}(\sigma_{attr3 \geq 6}(table1))$$

table1		
attr1	attr2	attr3
1	2	5
3	4	6
1	2	7
1	2	7

```
SELECT attr2
FROM table1
WHERE attr3 >= 6;
```

Cross Product

$table1 \times table2$

table1		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

$table1 \times table2$			
A	B	B	D
1	2	2	3
1	2	4	5
1	2	4	5
1	2	2	3
1	2	4	5
1	2	4	5

```
SELECT *
FROM   table1
CROSS JOIN table2;
```

Cross Product

table1		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

$table1 \times table2$			
A	B	B	D
1	2	2	3
1	2	4	5
1	2	4	5
1	2	2	3
1	2	4	5
1	2	4	5

```
SELECT *
FROM   table1 , table2 ;
```

Cross Product - Projecting out duplicate columns

$$\pi_{A,B,D}(table1 \times table2)$$

table1		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

table1 \times table2		
A	B	D
1	2	3
1	2	5
1	2	5
1	2	3
1	2	5
1	2	5

```
SELECT A, table1.B, D
FROM   table1
CROSS JOIN table2;
```

Natural Join

table1		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

table1 ⋈ table2			
A	B	B	D
1	2	2	3
1	2	2	3

- In the relational operator, the duplicate column gets projected out
- In SQL, **SELECT** clause decides for the columns to be retrieved
- The * specifies retrieving all the columns

```

SELECT *
FROM table1
JOIN table2
ON table1.B = table2.B;

```

Natural Join - Projecting out Duplicate Columns

$$\pi_{A, \text{table1.B}, D}(\text{table1} \bowtie \text{table2})$$

table		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

```

SELECT A, table1.B, C
FROM   table1
JOIN   table2
ON     table1.B = table2.B;

```


Natural Join - Projecting out Duplicate Columns

table1		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

table1 ⋈ table2		
A	B	D
1	2	3
1	2	3

```

SELECT A, table2.B, C
FROM table1
JOIN table2
ON table1.B = table2.B;

```

Theta Join

table1		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

table1 ⋈ table2			
A	B	B	D
1	2	4	5
1	2	4	5
1	2	4	5
1	2	4	5

```

SELECT A, table1.B, C
FROM   table1
JOIN   table2
WHERE  table1.B < table2.B;

```

Natural Join AND Theta Join

table1 ⋈ *table2*
table1.B = table2.B & table1.A < table2.D

table		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

```

SELECT A, table1.B, C
FROM   table1
JOIN   table2
ON     table1.B = table2.B
WHERE  table1.A < table2.D;
```

Re-naming

$\rho(\text{RESULT}(A1, B1, B2, D1), \text{table1} \bowtie_{\text{table1.B}=\text{table2.B} \& \text{table1.A} < \text{table2.D}} \text{table2})$

table		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

```

SELECT  A AS A1, table1.B AS B1,
        table2.B AS B2, D AS D1
FROM    table1
JOIN    table2
ON      table1.B = table2.B
WHERE   table1.A < table2.D;

```

Re-naming

$\rho(\text{Result}(A1, B1, B2, D1), \text{table1}) \bowtie_{\text{table1.B}=\text{table2.B} \& \text{table1.A} < \text{table2.D}} \text{table2}$

table		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

Result			
A1	B1	B2	D1
1	2	2	3
1	2	2	3

```
CREATE TABLE Result (
  SELECT  A AS A1, table1.B AS B1,
          table2.B AS B2, D AS D1
  FROM    table1
  JOIN    table2
  ON      table1.B = table2.B
  WHERE   table1.A < table2.D;
);
```

Distinct

 $\pi_{attr1, attr2}(table1)$

table1		
attr1	attr2	attr3
1	2	5
3	4	6
1	2	7
1	2	7

table1	
attr1	attr2
1	2
3	4

```
SELECT DISTINCT attr1 , attr2
FROM      table1
```

Aggregation Operations - SUM

Example

table1	
A	B
1	2
3	4
1	2
1	2

SUM(B)
10

```
SELECT SUM(B)
FROM   table1;
```

Aggregation Operations - Average

Example

table1	
A	B
1	2
3	4
1	2
1	2

AVG(A)
1.5

```
SELECT AVG(A)
FROM   table1;
```


Aggregation Operations - MIN

Example

table1	
A	B
1	2
3	4
1	2
1	2

MIN(A)
1

```
SELECT MIN(A)
FROM   table1;
```

Aggregation Operations - MAX

Example

table1	
A	B
1	2
3	4
1	2
1	2

MAX(A)
3

```
SELECT MAX(A)
FROM   table1;
```

Aggregation Operations - COUNT

Example

table1	
A	B
1	2
3	4
1	2
1	2

COUNT(A)
4

```
SELECT COUNT(A)
FROM   table1;
```

Extended Projection

 $\pi_{A, B+C \rightarrow X}(table1)$

table1		
A	B	C
0	1	2
0	1	2
3	4	5

A	X
0	3
0	3
3	9

```
SELECT A, (B + C) AS X
FROM   table1;
```

Extended Projection

 $\pi_{B-A \rightarrow X, C-B \rightarrow Y}(table1)$

table1		
A	B	C
0	1	2
0	1	2
3	4	5

X	Y
1	1
1	1
1	1

```
SELECT (B - A) AS X, (C - B) AS Y
FROM   table1;
```

Sorting

 $\tau_A(table1)$

table1		
A	B	C
3	4	5
1	1	2
7	1	2

table1		
A	B	C
1	1	2
3	4	5
7	1	2

```
SELECT A, B, C
FROM   table1;
ORDER BY A;
```

Right Outer Join

Right Outer Join

U		
A	B	C
1	2	3
4	5	6
7	8	9

V		
B	C	D
2	3	10
2	3	11
6	7	12

$U \overset{\circ}{\bowtie} V$ R					
A	B	C	B	C	D
1	2	3	2	3	10
1	2	3	2	3	11
\perp	\perp	\perp	6	7	12

```

SELECT *
FROM U
RIGHT OUTER JOIN V
ON U.B = V.B
AND U.C = V.C;

```

Left Outer Join

Left Outer Join

U		
A	B	C
1	2	3
4	5	6
7	8	9

V		
B	C	D
2	3	10
2	3	11
6	7	12

$U \overset{\circ}{\bowtie}_L V$					
A	B	C	B	C	D
1	2	3	2	3	10
1	2	3	2	3	11
4	5	6	⊥	⊥	⊥
7	8	9	⊥	⊥	⊥

```

SELECT *
FROM U
LEFT OUTER JOIN V
ON U.B = V.B
AND U.C = V.C;

```


Grouping

 $\gamma_{rating}(Sailors)$

Sailors			
sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

$\gamma_{rating}(Sailors)$			
sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
58	Rusty	10	35.0
74	Horatio	9	35.0
85	Art	3	25.5

Grouping

Group by rating

```
SELECT *
FROM   Sailors
GROUP BY rating;
```

Output

Sailors			
sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

$\gamma_{rating}(Sailors)$			
sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
58	Rusty	10	35.0
74	Horatio	9	35.0
85	Art	3	25.5

Grouping

Group by rating such that each group has at least two sailors

```
SELECT *
FROM   Sailors
GROUP BY rating
HAVING COUNT(rating) > 1;
```

Output

Sailors			
sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5

$\gamma_{rating}(Sailors)$			
sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.0
85	Art	3	25.5

Grouping

Group by rating such that each group has at least two sailors where sailor age ≥ 30

```
SELECT *
FROM   Sailors
WHERE  age >= 30
GROUP BY rating
HAVING COUNT(rating) > 1;
```

Output

Sailors			
sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0

$\gamma_{rating}(Sailors)$			
sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.0

All six clauses of SELECT

Group by rating such that each group has at least two sailors where sailor age ≥ 20 sort by sailor names

```
SELECT *
FROM   Sailors
WHERE  age >= 20
GROUP BY rating
HAVING COUNT(rating) > 1
ORDER BY sname;
```

Output

Sailors			
sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0

$\gamma_{rating}(Sailors)$			
sid	sname	rating	age
85	Art	3	25.5
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.0

Set Operator - Union

table1 \cup table2

table1		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

table1 \cup table2	
A	B
1	2
2	3
4	5

```
(SELECT      *
FROM      table1)
```

```
UNION
```

```
(SELECT      *
FROM      table2);
```

Set Operator - Intersection

table1 \cap table2

table1		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

table1 \cap table2	
A	B

```

SELECT      *
FROM        table1
WHERE (a, b)

IN

(SELECT      *
FROM        table2);

```

Set Operator - Difference

table1 – table2

table1		table2	
A	B	B	D
1	2	2	3
1	2	4	5
		4	5

table1 – table2	
A	B
1	2
2	3
4	5

```

SELECT  *
FROM    table1
WHERE   (a, b)

NOT IN

(SELECT  *
FROM    table2);

```